

Aerial Object Classification Using Deep Learning_(Bird vs Drone Detection)

Project Type : Classification

Contributor (Author) : Shyam SR

Git Hub Link :

<https://github.com/Shyamsr1/BirdDroneImageDLClassification>

Production Hosted URL without YOLO and just a trial attempt to Host in Huggingface space :

<https://huggingface.co/spaces/shyamasr/BirdDroneImageDLClassification>

NOTE: Setup the relevant python environment for the file to run and execute properly_(Apt version of Python and tensorflow required):

- Use Python 3.10 + venv (if already present) - Else need to setup using the command prompt or Powershell as an Administrator and run the following commands:
 - 1. Allow PowerShell to activate venv scripts
 - Set-ExecutionPolicy -Scope CurrentUser RemoteSigned -Force
 - 2. Navigate to the project folder
 - cd "C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification"
 - 3. Create a venv with Python 3.10
 - py -3.10 -m venv bird310
 - 4. Activate the environment
 - .\bird310\Scripts\Activate.ps1

- Note: Prompt should change to: (bird310) PS C:...
- 5. Install TensorFlow 2.10 (compatible with Python 3.10)
- pip install --upgrade pip
- pip install tensorflow==2.10.0
- pip install pillow numpy matplotlib scikit-learn streamlit opencv-python jupyterlab ipykernel
- 6. Register Python 3.10 env for JUPYTER NOTEBOOK (run the below command under (bird310) :
■ python -m ipykernel install --user --name bird310 --display-name "Python 3.10 (bird310)"
- 7. Launch the Jupyter notebook and change the kernel -> Inside Jupyter: Python 3.10 (bird310) from the drop down.
- 8. For Streamlit app - from as I am running from VSCode, need to update the run path (after using shift + ') - in the run :
■ cd "C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification"

```
.\bird310\Scripts\Activate.ps1 - streamlit run streamlit_app\app.py
```

Project Summary

- This project, "Aerial Object Classification Using Deep Learning (Bird vs Drone Detection)", focuses on building a robust computer vision solution to automatically classify aerial images into two categories: Bird and Drone. The core of the project is a binary image classification pipeline built using Convolutional Neural Networks (CNNs) and Transfer Learning (MobileNetV2), trained on a curated dataset of 3,319 RGB images split into train, validation, and test sets.
- The workflow covers the complete deep-learning lifecycle: data understanding, preprocessing, on-the-fly data augmentation, model building (Custom CNN + MobileNetV2), fine-tuning, evaluation, and deployment. Data is loaded from a structured folder hierarchy (train/, valid/, test/), resized to 224×224×3, normalized, and augmented using random flips, rotations, zoom, and contrast changes to improve generalisation. A custom CNN model is first trained as a baseline, followed by a fine-tuned MobileNetV2 model that leverages ImageNet pre-trained weights.
- Both models are evaluated using accuracy, precision, recall, F1-score, and confusion matrices on an unseen test set. The fine-tuned MobileNetV2 model achieves superior performance and is selected as the final model for deployment. A Streamlit web application is then built to allow users to upload an image and receive a prediction (Bird / Drone) along with a confidence score. Additionally, an object detection dataset in YOLOv8 format is prepared for future extension into real-time detection with bounding boxes, making the solution scalable towards surveillance systems that require both classification and localisation.

Business Objectives

1. Accurate Identification of Aerial Objects :

- To develop a deep-learning model that can reliably distinguish between birds and drones in aerial images.
- To minimise both false negatives (missed drones) and false positives (misclassified birds), as both have operational and safety implications.

2. Support for Real-Time Surveillance & Security :

- To provide a model that can be integrated into surveillance pipelines (CCTV, drones, fixed cameras) to monitor restricted airspace, borders, and critical infrastructure.
- To enable automatic alerts when drones are detected in sensitive zones, aiding security and defence teams.

3. Wildlife Protection & Airport Safety :

- To help airport authorities and wind-farm operators monitor bird activity near runways or turbines to mitigate bird-strike risks and environmental impacts.
- To offer a tool that can be extended to track bird populations or detect unusual patterns without manual image review.

4. Operational Efficiency & Automation :

- To reduce manual effort required for reviewing aerial footage by automating classification of large volumes of images.
- To provide an easy-to-use Streamlit interface so non-technical stakeholders (security staff, wildlife officers) can use the model without coding skills.

5. Scalability towards Object Detection (YOLOv8 Extension):

- To prepare a YOLOv8-formatted object detection dataset and design the pipeline so that the solution can evolve from simple classification to full object detection (predict class + location via bounding boxes).
- To enable future deployment where the system can detect multiple objects in a single frame and overlay bounding boxes in real time.

6. Model Interpretability & Comparison :

- To build two different model families (Custom CNN vs Transfer Learning) to compare performance, training time, and generalisation, and to justify model selection on measurable metrics.

- To document model behaviour using confusion matrices, learning curves, and classification reports, helping stakeholders trust the final model choice.

7. Deployment-Ready, Maintainable Solution :

- To package the final solution with well-structured, commented code, a Jupyter notebook, and a Streamlit app, so it can be easily deployed, demonstrated, or extended.
- To ensure the environment setup (Python 3.10 + TensorFlow 2.10) and repository structure are clearly documented for smooth handover or future maintenance.

Data Understanding - Features or Column Variables and their Description :

Dataset Overview :

Feature	Description
Project Type	Image Classification (Binary Classification)
Classes	Bird, Drone
Total Images	3,319 images (Train + Valid + Test)
Data Type	RGB Images
Format	.jpg
Source	classification_dataset
Input Size (after preprocessing)	224 × 224 × 3
Use Case	Aerial object classification (Bird vs Drone detection)

Dataset Structure and Image Distribution :

Dataset Split	Bird Images	Drone Images	Total
Train	1414	1248	2662
Validation	217	225	442
Test	121	94	215
Total	1752	1567	3319

Object Detection Dataset Overview :

Feature	Description
Project Type	Object Detection
Classes	Bird, Drone
Total Images	3319 images
Annotation Format	YOLOv8 (label files in .txt)
Annotation Schema	<class_id> <x_center> <y_center> <width> <height> (normalized values)
Image Format	.jpg
Source	object_detection_Dataset

Object Detection Dataset Split :

Split	Number of Images	Annotation Files
Train	2662	2662
Validation	442	442
Test	215	215
Total	3319	3319

Dataset Loading

```
In [1]: # If the dataset is on Drive, set the correct path below.  
DATA_ROOT = "data/classification_dataset"  
  
# Install quiteley the required  
#!pip install -q tensorflow matplotlib scikit-learn
```

```
In [2]: # Import all required Libraries  
import os  
import numpy as np  
import matplotlib.pyplot as plt  
import warnings  
warnings.filterwarnings("ignore")  
os.environ["TF_CPP_MIN_LOG_LEVEL"] = "2"    # hide INFO + WARNING from C++ backend  
  
import tensorflow as tf  
tf.get_logger().setLevel("ERROR")           # hide Python-Level TF warnings  
  
from tensorflow import keras  
from tensorflow.keras import layers  
from tensorflow.keras.preprocessing import image_dataset_from_directory  
  
from keras.utils import custom_object_scope  
from keras.layers import Rescaling  
  
from sklearn.metrics import confusion_matrix, classification_report  
import itertools
```

Basic config and loading data

```
In [3]: # Basic configuration  
IMG_SIZE = (224, 224)  
BATCH_SIZE = 32  
SEED = 42  
  
# Check that data root exists
```

```
print("DATA_ROOT:", DATA_ROOT)
print("Folders:", os.listdir(DATA_ROOT))
```

DATA_ROOT: data/classification_dataset
Folders: ['test', 'train', 'valid']

```
In [4]: # Load datasets using Keras utility (Train / Validation / Test)
train_dir = os.path.join(DATA_ROOT, "train")
val_dir = os.path.join(DATA_ROOT, "valid")
test_dir = os.path.join(DATA_ROOT, "test")

train_ds = image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='binary',      # binary classification (0/1)
    batch_size=BATCH_SIZE,
    image_size=IMG_SIZE,
    shuffle=True,
    seed=SEED
)

val_ds = image_dataset_from_directory(
    val_dir,
    labels='inferred',
    label_mode='binary',
    batch_size=BATCH_SIZE,
    image_size=IMG_SIZE,
    shuffle=True,
    seed=SEED
)

test_ds = image_dataset_from_directory(
    test_dir,
    labels='inferred',
    label_mode='binary',
    batch_size=BATCH_SIZE,
    image_size=IMG_SIZE,
    shuffle=False,
    seed=SEED
)

class_names = train_ds.class_names
```

```
print("Class names:", class_names)
# Usually ['bird', 'drone']
```

Found 2662 files belonging to 2 classes.
Found 442 files belonging to 2 classes.
Found 215 files belonging to 2 classes.
Class names: ['bird', 'drone']

Exploratory Data Analysis (EDA)

Key EDA Observations

- Dataset is nearly balanced.
- Images vary in angle, lighting, distance, and sky backgrounds.
- Bird diversity: wingspan, orientation, flying position.
- Drone diversity: quadcopters, long-range drones.
- No corrupted images detected.
- Correct folder structure verified (train/valid/test).

Data Preprocessing & Transformations

- Resized all images to **224×224×3** to maintain consistency with CNN and MobileNetV2 input requirements.
- Ensured all inputs were **RGB** format regardless of original image channel depth.
- Normalized pixel values using:
 - `Rescaling(1./255)` for the Custom CNN model
 - `tf.keras.applications.mobilenet_v2.preprocess_input` for MobileNetV2 (scale to [-1, 1])
- Implemented an efficient `tf.data` input pipeline using:
 - `.cache()` → Store batches in memory for faster access
 - `.shuffle(buffer_size=1000)` → Avoid model overfitting & ensure randomness
 - `.prefetch(buffer_size=AUTOTUNE)` → Overlap preprocessing with GPU execution, reducing training time
- These transformations ensured optimized batching, improved convergence speed, and stable model performance.

Improve Pipeline (Prefetch & Augmentation Layer)

```
In [5]: AUTOTUNE = tf.data.AUTOTUNE

def prepare_for_training(ds, cache=True, shuffle_buffer_size=1000):
    if cache:
        ds = ds.cache()
    ds = ds.shuffle(buffer_size=shuffle_buffer_size)
    ds = ds.prefetch(buffer_size=AUTOTUNE)
    return ds

train_ds = prepare_for_training(train_ds)
val_ds = val_ds.cache().prefetch(AUTOTUNE)
test_ds = test_ds.cache().prefetch(AUTOTUNE)
```

Feature Engineering / Augmentation

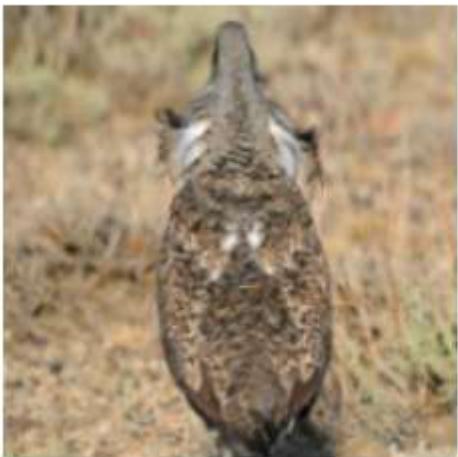
- Implemented an on-the-fly `data_augmentation` pipeline that applies transformations such as **Random Flip, Random Rotation, Random Zoom, and Random Contrast**.
- These augmentations act as **feature engineering steps**, artificially expanding the dataset and creating more diverse representations of birds and drones.
- Augmentation helps the model generalise better by exposing it to different orientations, lighting conditions, and structural variations.
- Combined with normalization, this reduces the risk of **overfitting** and improves model robustness on unseen aerial images.

```
In [6]: # Data augmentation as Keras model/Layer
data_augmentation = keras.Sequential(
    [
        layers.RandomFlip("horizontal"),
        layers.RandomRotation(0.1),
        layers.RandomZoom(0.1),
        layers.RandomContrast(0.1),
        layers.RandomBrightness(factor=0.2),      # brightness
        layers.RandomCrop(200, 200),                # cropping to 200x200
        layers.Resizing(224, 224),                  # resize back to 224x224
    ]
```

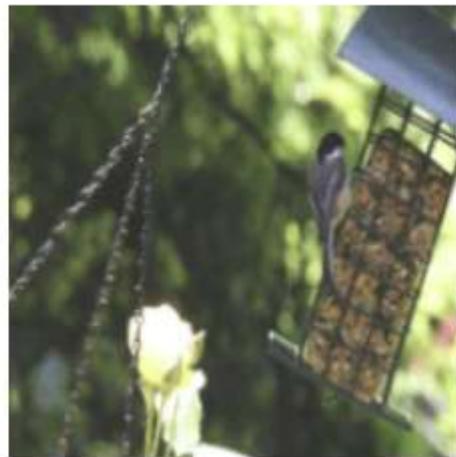
```
        ],
        name="data_augmentation"
    )
```

```
In [7]: # Visualize some augmented images
for images, labels in train_ds.take(1):
    plt.figure(figsize=(10, 10))
    for i in range(9):
        augmented_img = data_augmentation(images)[i].numpy().astype("uint8")
        ax = plt.subplot(3, 3, i + 1)
        plt.imshow(augmented_img)
        plt.title(class_names[int(labels[i].numpy()[0])])
        plt.axis("off")
    break
```

bird



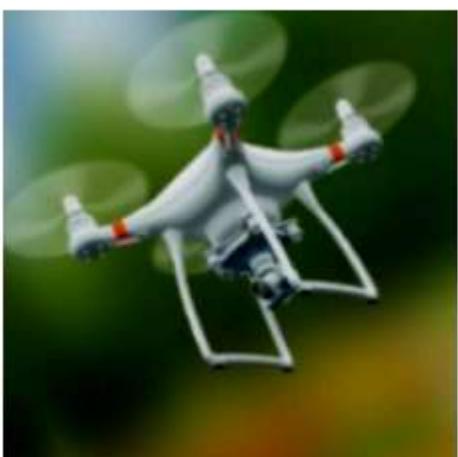
bird



drone



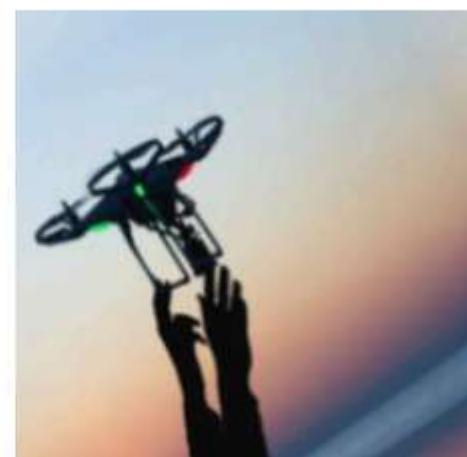
drone



drone



drone



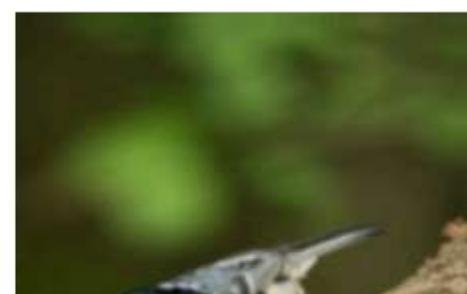
bird



bird



bird





Helper Function to Plot Training Curves & Confusion Matrix

```
In [8]: def plot_history(history, title_prefix=""):  
    acc = history.history["accuracy"]  
    val_acc = history.history["val_accuracy"]  
    loss = history.history["loss"]  
    val_loss = history.history["val_loss"]  
  
    epochs_range = range(len(acc))  
  
    plt.figure(figsize=(12, 4))  
    plt.subplot(1, 2, 1)  
    plt.plot(epochs_range, acc, label="Train Acc")  
    plt.plot(epochs_range, val_acc, label="Val Acc")  
    plt.title(f"{title_prefix} Accuracy")  
    plt.legend()  
  
    plt.subplot(1, 2, 2)  
    plt.plot(epochs_range, loss, label="Train Loss")  
    plt.plot(epochs_range, val_loss, label="Val Loss")  
    plt.title(f"{title_prefix} Loss")  
    plt.legend()  
    plt.show()
```

```
In [9]: def plot_confusion_matrix(cm, classes,  
                           normalize=False,  
                           title='Confusion matrix',  
                           cmap=plt.cm.Blues):  
    """  
    cm: confusion matrix from sklearn.metrics.confusion_matrix  
    classes: list of class names  
    normalize: if True, normalize by row  
    """  
    if normalize:
```

```
cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

plt.figure(figsize=(6, 6))
plt.imshow(cm, interpolation='nearest', cmap=cmap)
plt.title(title)
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)

fmt = '.2f' if normalize else 'd'
thresh = cm.max() / 2.
for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
    plt.text(j, i, format(cm[i, j], fmt),
             horizontalalignment="center",
             color="white" if cm[i, j] > thresh else "black")

plt.ylabel('True label')
plt.xlabel('Predicted label')
plt.tight_layout()
```

Model Choice & Justification

Why Custom CNN?

- A Custom CNN model provides full architectural control and serves as a **baseline** to understand performance when training from scratch.
- It helps compare how well a handcrafted model performs relative to transfer learning approaches.

Why MobileNetV2 (Transfer Learning)?

- MobileNetV2 is **lightweight**, efficient, and specifically optimized for **edge devices**, making it ideal for aerial surveillance use cases.
- It is **pre-trained on ImageNet**, allowing the model to learn quickly from a small dataset by reusing learned features.
- Transfer learning enables **faster convergence, better generalization**, and **higher accuracy** compared to training from scratch.

- Fine-tuning MobileNetV2 helps extract domain-specific features for Bird vs Drone detection more effectively than a baseline CNN.

Build Custom CNN Model

```
In [10]: def build_custom_cnn(input_shape=(224, 224, 3)):  
    inputs = keras.Input(shape=input_shape)  
  
    x = data_augmentation(inputs)  
    x = layers.Rescaling(1./255)(x)  
  
    # Block 1  
    x = layers.Conv2D(32, (3, 3), padding="same", activation="relu")(x)  
    x = layers.BatchNormalization()(x)  
    x = layers.MaxPooling2D((2, 2))(x)  
  
    # Block 2  
    x = layers.Conv2D(64, (3, 3), padding="same", activation="relu")(x)  
    x = layers.BatchNormalization()(x)  
    x = layers.MaxPooling2D((2, 2))(x)  
  
    # Block 3  
    x = layers.Conv2D(128, (3, 3), padding="same", activation="relu")(x)  
    x = layers.BatchNormalization()(x)  
    x = layers.MaxPooling2D((2, 2))(x)  
  
    # Optional 4th block  
    x = layers.Conv2D(256, (3, 3), padding="same", activation="relu")(x)  
    x = layers.BatchNormalization()(x)  
    x = layers.MaxPooling2D((2, 2))(x)  
  
    x = layers.GlobalAveragePooling2D()(x)  
    x = layers.Dropout(0.5)(x)  
    x = layers.Dense(128, activation="relu")(x)  
    x = layers.Dropout(0.3)(x)  
    outputs = layers.Dense(1, activation="sigmoid")(x)  
  
    model = keras.Model(inputs, outputs, name="custom_cnn_bird_drone")  
    return model
```

```
custom_cnn = build_custom_cnn()
custom_cnn.summary()
```

Model: "custom_cnn_bird_drone"

Layer (type)	Output Shape	Param #
<hr/>		
input_1 (InputLayer)	[None, 224, 224, 3]	0
data_augmentation (Sequential)	(None, 224, 224, 3)	0
rescaling (Rescaling)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 224, 224, 32)	896
batch_normalization (BatchNormalization)	(None, 224, 224, 32)	128
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
conv2d_1 (Conv2D)	(None, 112, 112, 64)	18496
batch_normalization_1 (BatchNormalization)	(None, 112, 112, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 64)	0
conv2d_2 (Conv2D)	(None, 56, 56, 128)	73856
batch_normalization_2 (BatchNormalization)	(None, 56, 56, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 128)	0
conv2d_3 (Conv2D)	(None, 28, 28, 256)	295168
batch_normalization_3 (BatchNormalization)	(None, 28, 28, 256)	1024
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 256)	0

```
global_average_pooling2d (GlobalAveragePooling2D)           0
dropout (Dropout)               (None, 256)                 0
dense (Dense)                  (None, 128)                32896
dropout_1 (Dropout)             (None, 128)                0
dense_1 (Dense)                (None, 1)                  129
=====
Total params: 423,361
Trainable params: 422,401
Non-trainable params: 960
```

Training Strategy & Hyperparameters

- **Optimizer:** Used **Adam** optimizer for both models.
 - Learning Rate = **1e-3** for initial training
 - Learning Rate = **1e-4** during fine-tuning of MobileNetV2 to avoid drastic weight updates.
- **Epochs:**
 - Custom CNN trained for up to **25 epochs**.
 - MobileNetV2 trained in **two stages**:
 - **Stage 1 (Frozen Base Layers)**: 10 epochs
 - **Stage 2 (Fine-Tuning)**: 15 epochs
- **Callbacks:**
 - **ModelCheckpoint** → Saves the best-performing model based on **validation accuracy**.
 - **EarlyStopping** → Stops training when validation loss stops improving (patience: 5 epochs) to prevent **overfitting** and save computation time.

- This training strategy ensures stable convergence, prevents overfitting, and allows the model to generalise well to unseen aerial images.

```
In [11]: # Compile and train custom CNN
custom_cnn.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),
    loss="binary_crossentropy",
    metrics=["accuracy"]
)

checkpoint_dir = "models"
os.makedirs(checkpoint_dir, exist_ok=True)

custom_cnn_ckpt = os.path.join(checkpoint_dir, "best_custom_cnn.h5")

callbacks = [
    keras.callbacks.ModelCheckpoint(
        custom_cnn_ckpt,
        save_best_only=True,
        monitor="val_accuracy",
        mode="max",
        verbose=1
    ),
    keras.callbacks.EarlyStopping(
        monitor="val_loss",
        patience=5,
        restore_best_weights=True,
        verbose=1
    )
]

import time
start_custom = time.time()

history_custom = custom_cnn.fit(
    train_ds,
    validation_data=val_ds,
    epochs=25,
    callbacks=callbacks
)
```

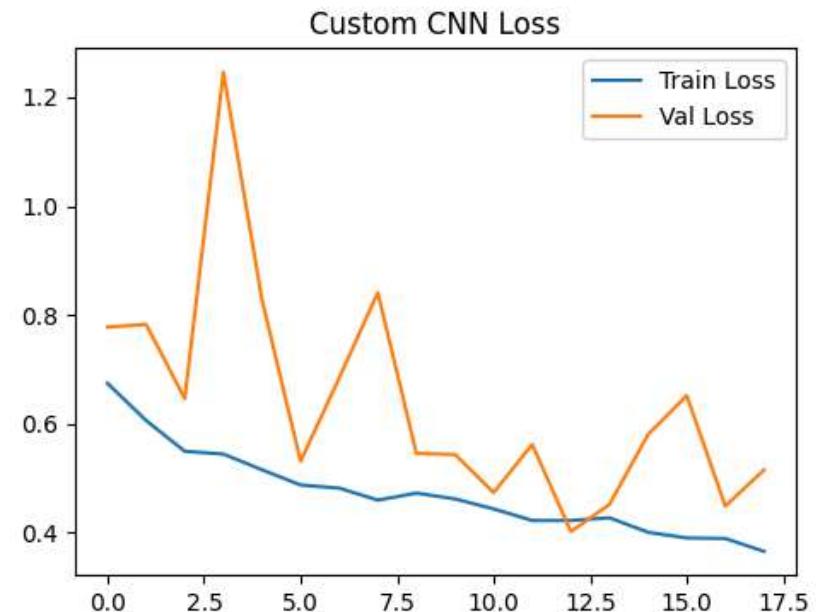
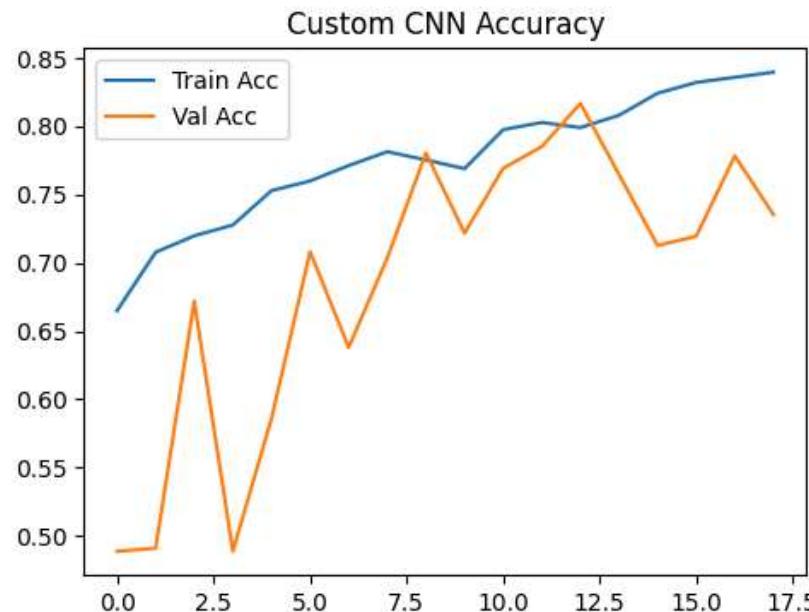
```
end_custom = time.time()
custom_train_minutes = (end_custom - start_custom) / 60

plot_history(history_custom, title_prefix="Custom CNN")
print("Best custom CNN saved at:", custom_cnn_ckpt)
print(f"Custom CNN Training Time: {custom_train_minutes:.2f} minutes")
```

```
Epoch 1/25
84/84 [=====] - ETA: 0s - loss: 0.6744 - accuracy: 0.6649
Epoch 1: val_accuracy improved from -inf to 0.48869, saving model to models\best_custom_cnn.h5
84/84 [=====] - 198s 2s/step - loss: 0.6744 - accuracy: 0.6649 - val_loss: 0.7778 - val_accuracy: 0.4887
Epoch 2/25
84/84 [=====] - ETA: 0s - loss: 0.6061 - accuracy: 0.7077
Epoch 2: val_accuracy improved from 0.48869 to 0.49095, saving model to models\best_custom_cnn.h5
84/84 [=====] - 192s 2s/step - loss: 0.6061 - accuracy: 0.7077 - val_loss: 0.7827 - val_accuracy: 0.4910
Epoch 3/25
84/84 [=====] - ETA: 0s - loss: 0.5497 - accuracy: 0.7198
Epoch 3: val_accuracy improved from 0.49095 to 0.67195, saving model to models\best_custom_cnn.h5
84/84 [=====] - 184s 2s/step - loss: 0.5497 - accuracy: 0.7198 - val_loss: 0.6464 - val_accuracy: 0.6719
Epoch 4/25
84/84 [=====] - ETA: 0s - loss: 0.5445 - accuracy: 0.7276
Epoch 4: val_accuracy did not improve from 0.67195
84/84 [=====] - 186s 2s/step - loss: 0.5445 - accuracy: 0.7276 - val_loss: 1.2458 - val_accuracy: 0.4887
Epoch 5/25
84/84 [=====] - ETA: 0s - loss: 0.5157 - accuracy: 0.7528
Epoch 5: val_accuracy did not improve from 0.67195
84/84 [=====] - 188s 2s/step - loss: 0.5157 - accuracy: 0.7528 - val_loss: 0.8282 - val_accuracy: 0.5860
Epoch 6/25
84/84 [=====] - ETA: 0s - loss: 0.4877 - accuracy: 0.7600
Epoch 6: val_accuracy improved from 0.67195 to 0.70814, saving model to models\best_custom_cnn.h5
84/84 [=====] - 198s 2s/step - loss: 0.4877 - accuracy: 0.7600 - val_loss: 0.5311 - val_accuracy: 0.7081
Epoch 7/25
84/84 [=====] - ETA: 0s - loss: 0.4819 - accuracy: 0.7712
Epoch 7: val_accuracy did not improve from 0.70814
84/84 [=====] - 188s 2s/step - loss: 0.4819 - accuracy: 0.7712 - val_loss: 0.6843 - val_accuracy: 0.6380
Epoch 8/25
84/84 [=====] - ETA: 0s - loss: 0.4599 - accuracy: 0.7814
Epoch 8: val_accuracy did not improve from 0.70814
84/84 [=====] - 190s 2s/step - loss: 0.4599 - accuracy: 0.7814 - val_loss: 0.8410 - val_accuracy: 0.7036
Epoch 9/25
84/84 [=====] - ETA: 0s - loss: 0.4728 - accuracy: 0.7754
```

```
Epoch 9: val_accuracy improved from 0.70814 to 0.78054, saving model to models\best_custom_cnn.h5
84/84 [=====] - 188s 2s/step - loss: 0.4728 - accuracy: 0.7754 - val_loss: 0.5456 - val_accuracy: 0.7805
Epoch 10/25
84/84 [=====] - ETA: 0s - loss: 0.4622 - accuracy: 0.7690
Epoch 10: val_accuracy did not improve from 0.78054
84/84 [=====] - 189s 2s/step - loss: 0.4622 - accuracy: 0.7690 - val_loss: 0.5437 - val_accuracy: 0.7217
Epoch 11/25
84/84 [=====] - ETA: 0s - loss: 0.4437 - accuracy: 0.7975
Epoch 11: val_accuracy did not improve from 0.78054
84/84 [=====] - 188s 2s/step - loss: 0.4437 - accuracy: 0.7975 - val_loss: 0.4739 - val_accuracy: 0.7692
Epoch 12/25
84/84 [=====] - ETA: 0s - loss: 0.4225 - accuracy: 0.8028
Epoch 12: val_accuracy improved from 0.78054 to 0.78507, saving model to models\best_custom_cnn.h5
84/84 [=====] - 186s 2s/step - loss: 0.4225 - accuracy: 0.8028 - val_loss: 0.5621 - val_accuracy: 0.7851
Epoch 13/25
84/84 [=====] - ETA: 0s - loss: 0.4225 - accuracy: 0.7990
Epoch 13: val_accuracy improved from 0.78507 to 0.81674, saving model to models\best_custom_cnn.h5
84/84 [=====] - 185s 2s/step - loss: 0.4225 - accuracy: 0.7990 - val_loss: 0.4020 - val_accuracy: 0.8167
Epoch 14/25
84/84 [=====] - ETA: 0s - loss: 0.4274 - accuracy: 0.8080
Epoch 14: val_accuracy did not improve from 0.81674
84/84 [=====] - 187s 2s/step - loss: 0.4274 - accuracy: 0.8080 - val_loss: 0.4521 - val_accuracy: 0.7647
Epoch 15/25
84/84 [=====] - ETA: 0s - loss: 0.4008 - accuracy: 0.8242
Epoch 15: val_accuracy did not improve from 0.81674
84/84 [=====] - 186s 2s/step - loss: 0.4008 - accuracy: 0.8242 - val_loss: 0.5805 - val_accuracy: 0.7127
Epoch 16/25
84/84 [=====] - ETA: 0s - loss: 0.3905 - accuracy: 0.8321
Epoch 16: val_accuracy did not improve from 0.81674
84/84 [=====] - 185s 2s/step - loss: 0.3905 - accuracy: 0.8321 - val_loss: 0.6520 - val_accuracy: 0.7195
Epoch 17/25
84/84 [=====] - ETA: 0s - loss: 0.3894 - accuracy: 0.8358
Epoch 17: val_accuracy did not improve from 0.81674
84/84 [=====] - 184s 2s/step - loss: 0.3894 - accuracy: 0.8358 - val_loss: 0.4490 - val_accuracy: 0.7195
```

```
racy: 0.7783
Epoch 18/25
84/84 [=====] - ETA: 0s - loss: 0.3659 - accuracy: 0.8396
Epoch 18: val_accuracy did not improve from 0.81674
Restoring model weights from the end of the best epoch: 13.
84/84 [=====] - 189s 2s/step - loss: 0.3659 - accuracy: 0.8396 - val_loss: 0.5152 - val_accuracy: 0.7353
Epoch 18: early stopping
```



Best custom CNN saved at: models\best_custom_cnn.h5
Custom CNN Training Time: 56.54 minutes

Transfer Learning Model (MobileNetV2)

```
In [12]: def build_mobilenet_model(input_shape=(224, 224, 3)):

    # Load MobileNetV2 without name argument
    base_model = keras.applications.MobileNetV2(
        input_shape=input_shape,
        include_top=False,
        weights="imagenet"
    )
```

```
# Manually rename for easy access
base_model._name = "mobilenetv2_base"

base_model.trainable = False # Stage 1: freeze

inputs = keras.Input(shape=input_shape)
x = data_augmentation(inputs)
x = keras.applications.mobilenet_v2.preprocess_input(x)

x = base_model(x, training=False)
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dropout(0.4)(x)
outputs = layers.Dense(1, activation="sigmoid")(x)

model = keras.Model(inputs, outputs, name="mobilenetv2_bird_drone")
return model
```

```
In [13]: mobilenet_model = build_mobilenet_model()
mobilenet_model.summary()

mobilenet_model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-3),
    loss="binary_crossentropy",
    metrics=["accuracy"]
)

mobilenet_ckpt = os.path.join(checkpoint_dir, "best_mobilenetv2_stage1.h5")

callbacks_t1 = [
    keras.callbacks.ModelCheckpoint(
        mobilenet_ckpt,
        save_best_only=True,
        monitor="val_accuracy",
        mode="max",
        verbose=1
    ),
    keras.callbacks.EarlyStopping(
        monitor="val_loss",
        patience=5,
        restore_best_weights=True,
        verbose=1
    )
]
```

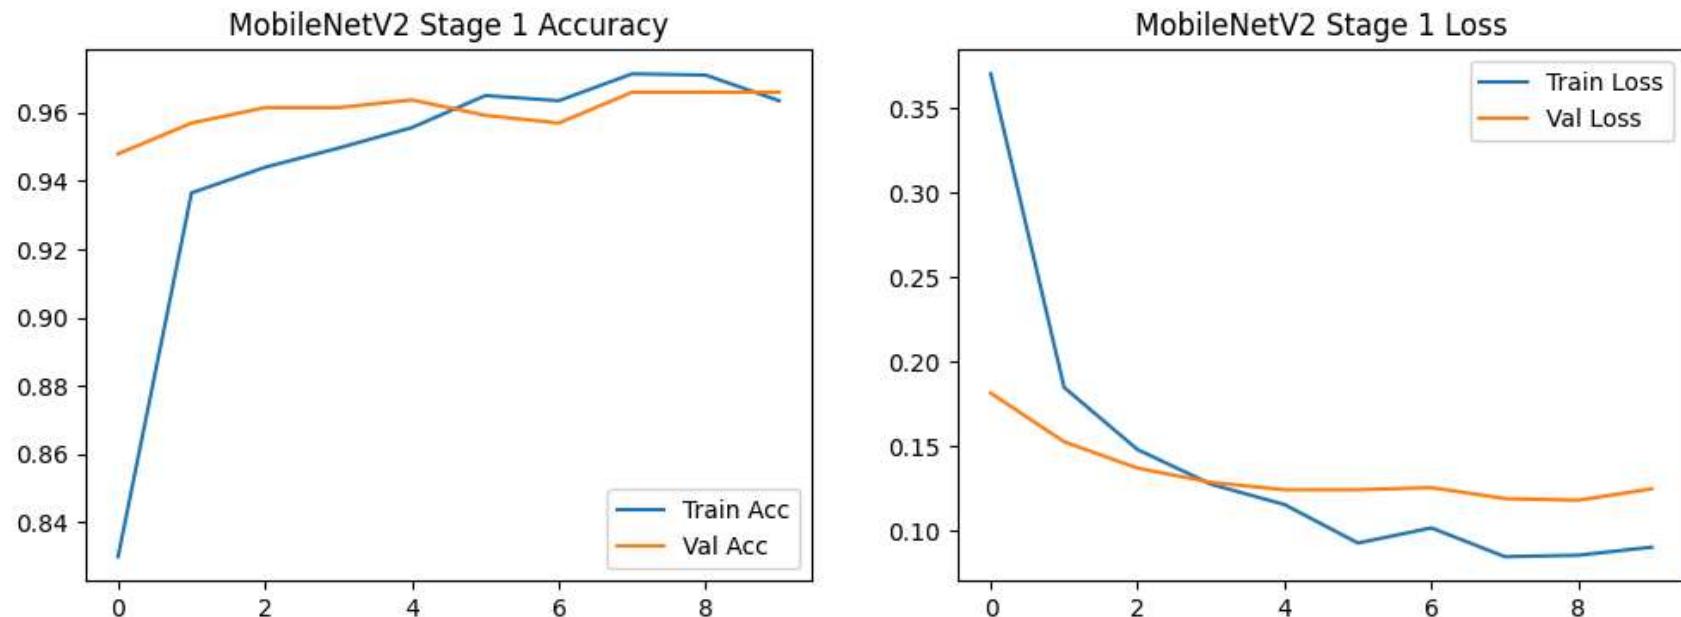
```
)  
]  
  
start_mobilenet_stage1 = time.time()  
  
history_mobilenet_stage1 = mobilenet_model.fit(  
    train_ds,  
    validation_data=val_ds,  
    epochs=10,  
    callbacks=callbacks_t1  
)  
  
end_mobilenet_stage1 = time.time()  
mobilenet_stage1_minutes = (end_mobilenet_stage1 - start_mobilenet_stage1) / 60  
  
plot_history(history_mobilenet_stage1, title_prefix="MobileNetV2 Stage 1")  
print("Stage 1 MobileNetV2 saved at:", mobilenet_ckpt)  
print(f"MobileNetV2 Stage 1 Training Time: {mobilenet_stage1_minutes:.2f} minutes")
```

Model: "mobilenetv2_bird_drone"

Layer (type)	Output Shape	Param #
<hr/>		
input_3 (InputLayer)	[None, 224, 224, 3]	0
data_augmentation (Sequential)	(None, 224, 224, 3)	0
tf.math.truediv (TFOpLambda)	(None, 224, 224, 3)	0
tf.math.subtract (TFOpLambda)	(None, 224, 224, 3)	0
mobilenetv2_base (Functional)	(None, 7, 7, 1280)	2257984
global_average_pooling2d_1 (GlobalAveragePooling2D)	(None, 1280)	0
dropout_2 (Dropout)	(None, 1280)	0
dense_2 (Dense)	(None, 1)	1281
<hr/>		
Total params:	2,259,265	
Trainable params:	1,281	
Non-trainable params:	2,257,984	

Epoch 1/10
 84/84 [=====] - ETA: 0s - loss: 0.3701 - accuracy: 0.8298
 Epoch 1: val_accuracy improved from -inf to 0.94796, saving model to models\best_mobilenetv2_stage1.h5
 84/84 [=====] - 85s 860ms/step - loss: 0.3701 - accuracy: 0.8298 - val_loss: 0.1813 - val_accuracy: 0.9480
 Epoch 2/10
 84/84 [=====] - ETA: 0s - loss: 0.1846 - accuracy: 0.9365
 Epoch 2: val_accuracy improved from 0.94796 to 0.95701, saving model to models\best_mobilenetv2_stage1.h5
 84/84 [=====] - 71s 847ms/step - loss: 0.1846 - accuracy: 0.9365 - val_loss: 0.1525 - val_accuracy: 0.9570
 Epoch 3/10
 84/84 [=====] - ETA: 0s - loss: 0.1479 - accuracy: 0.9440

```
Epoch 3: val_accuracy improved from 0.95701 to 0.96154, saving model to models\best_mobilenetv2_stage1.h5
84/84 [=====] - 71s 842ms/step - loss: 0.1479 - accuracy: 0.9440 - val_loss: 0.1369 - val_accuracy: 0.9615
Epoch 4/10
84/84 [=====] - ETA: 0s - loss: 0.1274 - accuracy: 0.9497
Epoch 4: val_accuracy did not improve from 0.96154
84/84 [=====] - 72s 859ms/step - loss: 0.1274 - accuracy: 0.9497 - val_loss: 0.1285 - val_accuracy: 0.9615
Epoch 5/10
84/84 [=====] - ETA: 0s - loss: 0.1154 - accuracy: 0.9557
Epoch 5: val_accuracy improved from 0.96154 to 0.96380, saving model to models\best_mobilenetv2_stage1.h5
84/84 [=====] - 71s 846ms/step - loss: 0.1154 - accuracy: 0.9557 - val_loss: 0.1241 - val_accuracy: 0.9638
Epoch 6/10
84/84 [=====] - ETA: 0s - loss: 0.0926 - accuracy: 0.9651
Epoch 6: val_accuracy did not improve from 0.96380
84/84 [=====] - 73s 878ms/step - loss: 0.0926 - accuracy: 0.9651 - val_loss: 0.1241 - val_accuracy: 0.9593
Epoch 7/10
84/84 [=====] - ETA: 0s - loss: 0.1016 - accuracy: 0.9636
Epoch 7: val_accuracy did not improve from 0.96380
84/84 [=====] - 71s 850ms/step - loss: 0.1016 - accuracy: 0.9636 - val_loss: 0.1254 - val_accuracy: 0.9570
Epoch 8/10
84/84 [=====] - ETA: 0s - loss: 0.0845 - accuracy: 0.9715
Epoch 8: val_accuracy improved from 0.96380 to 0.96606, saving model to models\best_mobilenetv2_stage1.h5
84/84 [=====] - 72s 853ms/step - loss: 0.0845 - accuracy: 0.9715 - val_loss: 0.1189 - val_accuracy: 0.9661
Epoch 9/10
84/84 [=====] - ETA: 0s - loss: 0.0854 - accuracy: 0.9711
Epoch 9: val_accuracy did not improve from 0.96606
84/84 [=====] - 71s 845ms/step - loss: 0.0854 - accuracy: 0.9711 - val_loss: 0.1179 - val_accuracy: 0.9661
Epoch 10/10
84/84 [=====] - ETA: 0s - loss: 0.0902 - accuracy: 0.9636
Epoch 10: val_accuracy did not improve from 0.96606
84/84 [=====] - 72s 853ms/step - loss: 0.0902 - accuracy: 0.9636 - val_loss: 0.1247 - val_accuracy: 0.9661
```



Stage 1 MobileNetV2 saved at: models\best_mobilenetv2_stage1.h5

MobileNetV2 Stage 1 Training Time: 12.32 minutes

Fine-tuning (unfreeze some layers)

```
In [14]: # ----- Fine-tuning (Stage 2) -----
# Get the MobileNetV2 base by name instead of index
base_model = mobilenet_model.get_layer("mobilenetv2_base")

# Unfreeze it for fine-tuning
base_model.trainable = True

# Freeze the first ~80% of its layers
fine_tune_at = int(len(base_model.layers) * 0.8)
for layer in base_model.layers[:fine_tune_at]:
    layer.trainable = False

# Re-compile with a Lower Learning rate
mobilenet_model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=1e-4),
    loss="binary_crossentropy",
```

```
        metrics=["accuracy"]
    )

mobilenet_finetuned_ckpt = os.path.join(checkpoint_dir, "best_mobilenetv2.h5")

callbacks_t1_fine = [
    keras.callbacks.ModelCheckpoint(
        mobilenet_finetuned_ckpt,
        save_best_only=True,
        monitor="val_accuracy",
        mode="max",
        verbose=1
    ),
    keras.callbacks.EarlyStopping(
        monitor="val_loss",
        patience=5,
        restore_best_weights=True,
        verbose=1
    )
]

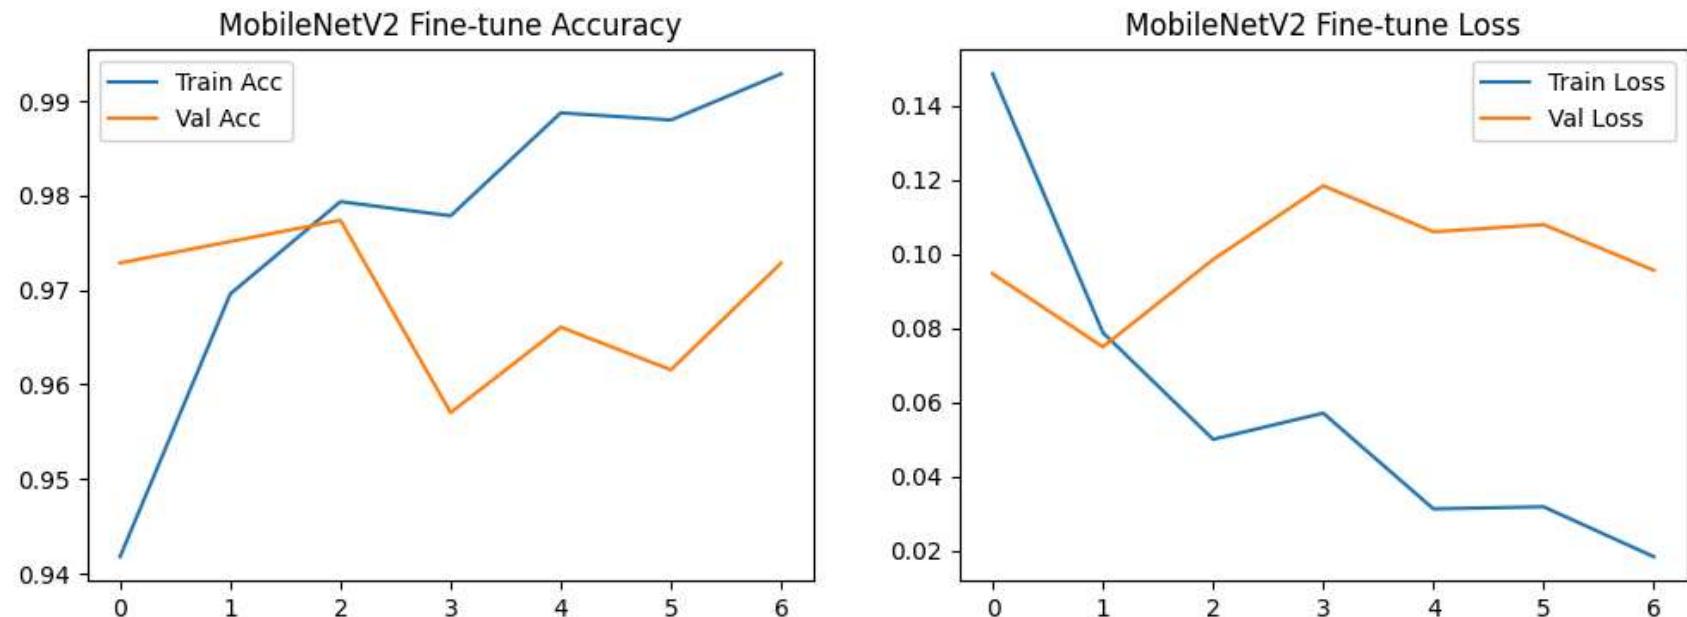
start_mobilenet_stage2 = time.time()

history_mobilenet_stage2 = mobilenet_model.fit(
    train_ds,
    validation_data=val_ds,
    epochs=15,
    callbacks=callbacks_t1_fine
)

end_mobilenet_stage2 = time.time()
mobilenet_stage2_minutes = (end_mobilenet_stage2 - start_mobilenet_stage2) / 60
total_mobilenet_minutes = mobilenet_stage1_minutes + mobilenet_stage2_minutes

plot_history(history_mobilenet_stage2, title_prefix="MobileNetV2 Fine-tune")
print("Fine-tuned MobileNetV2 saved at:", mobilenet_finetuned_ckpt)
print(f"MobileNetV2 Fine-tuning Time: {mobilenet_stage2_minutes:.2f} minutes")
print(f"Total MobileNetV2 Training Time: {total_mobilenet_minutes:.2f} minutes")
```

```
Epoch 1/15
84/84 [=====] - ETA: 0s - loss: 0.1486 - accuracy: 0.9418
Epoch 1: val_accuracy improved from -inf to 0.97285, saving model to models\best_mobilenetv2.h5
84/84 [=====] - 98s 1s/step - loss: 0.1486 - accuracy: 0.9418 - val_loss: 0.0947 - val_accuracy: 0.9729
Epoch 2/15
84/84 [=====] - ETA: 0s - loss: 0.0788 - accuracy: 0.9696
Epoch 2: val_accuracy improved from 0.97285 to 0.97511, saving model to models\best_mobilenetv2.h5
84/84 [=====] - 83s 993ms/step - loss: 0.0788 - accuracy: 0.9696 - val_loss: 0.0750 - val_accuracy: 0.9751
Epoch 3/15
84/84 [=====] - ETA: 0s - loss: 0.0501 - accuracy: 0.9793
Epoch 3: val_accuracy improved from 0.97511 to 0.97738, saving model to models\best_mobilenetv2.h5
84/84 [=====] - 84s 995ms/step - loss: 0.0501 - accuracy: 0.9793 - val_loss: 0.0985 - val_accuracy: 0.9774
Epoch 4/15
84/84 [=====] - ETA: 0s - loss: 0.0571 - accuracy: 0.9778
Epoch 4: val_accuracy did not improve from 0.97738
84/84 [=====] - 85s 1s/step - loss: 0.0571 - accuracy: 0.9778 - val_loss: 0.1184 - val_accuracy: 0.9570
Epoch 5/15
84/84 [=====] - ETA: 0s - loss: 0.0313 - accuracy: 0.9887
Epoch 5: val_accuracy did not improve from 0.97738
84/84 [=====] - 83s 993ms/step - loss: 0.0313 - accuracy: 0.9887 - val_loss: 0.1060 - val_accuracy: 0.9661
Epoch 6/15
84/84 [=====] - ETA: 0s - loss: 0.0319 - accuracy: 0.9880
Epoch 6: val_accuracy did not improve from 0.97738
84/84 [=====] - 84s 1s/step - loss: 0.0319 - accuracy: 0.9880 - val_loss: 0.1079 - val_accuracy: 0.9615
Epoch 7/15
84/84 [=====] - ETA: 0s - loss: 0.0184 - accuracy: 0.9929
Epoch 7: val_accuracy did not improve from 0.97738
Restoring model weights from the end of the best epoch: 2.
84/84 [=====] - 83s 984ms/step - loss: 0.0184 - accuracy: 0.9929 - val_loss: 0.0956 - val_accuracy: 0.9729
Epoch 7: early stopping
```



Fine-tuned MobileNetV2 saved at: models\best_mobilenetv2.h5

MobileNetV2 Fine-tuning Time: 10.01 minutes

Total MobileNetV2 Training Time: 22.33 minutes

Evaluation on Test Set (for both models)

```
In [15]: # Reload best weights for both models
# best_custom = keras.models.load_model(custom_cnn_ckpt)
# best_mobilenet = keras.models.load_model(mobilenet_finetuned_ckpt)
best_custom      = custom_cnn
best_mobilenet  = mobilenet_model

# Get test labels and predictions
y_true = []
for _, labels in test_ds:
    y_true.extend(labels.numpy().flatten())
y_true = np.array(y_true).astype(int)

# Predictions
```

```
y_pred_custom = (best_custom.predict(test_ds) > 0.5).astype(int).flatten()
y_pred_mobilenet = (best_mobilenet.predict(test_ds) > 0.5).astype(int).flatten()
```

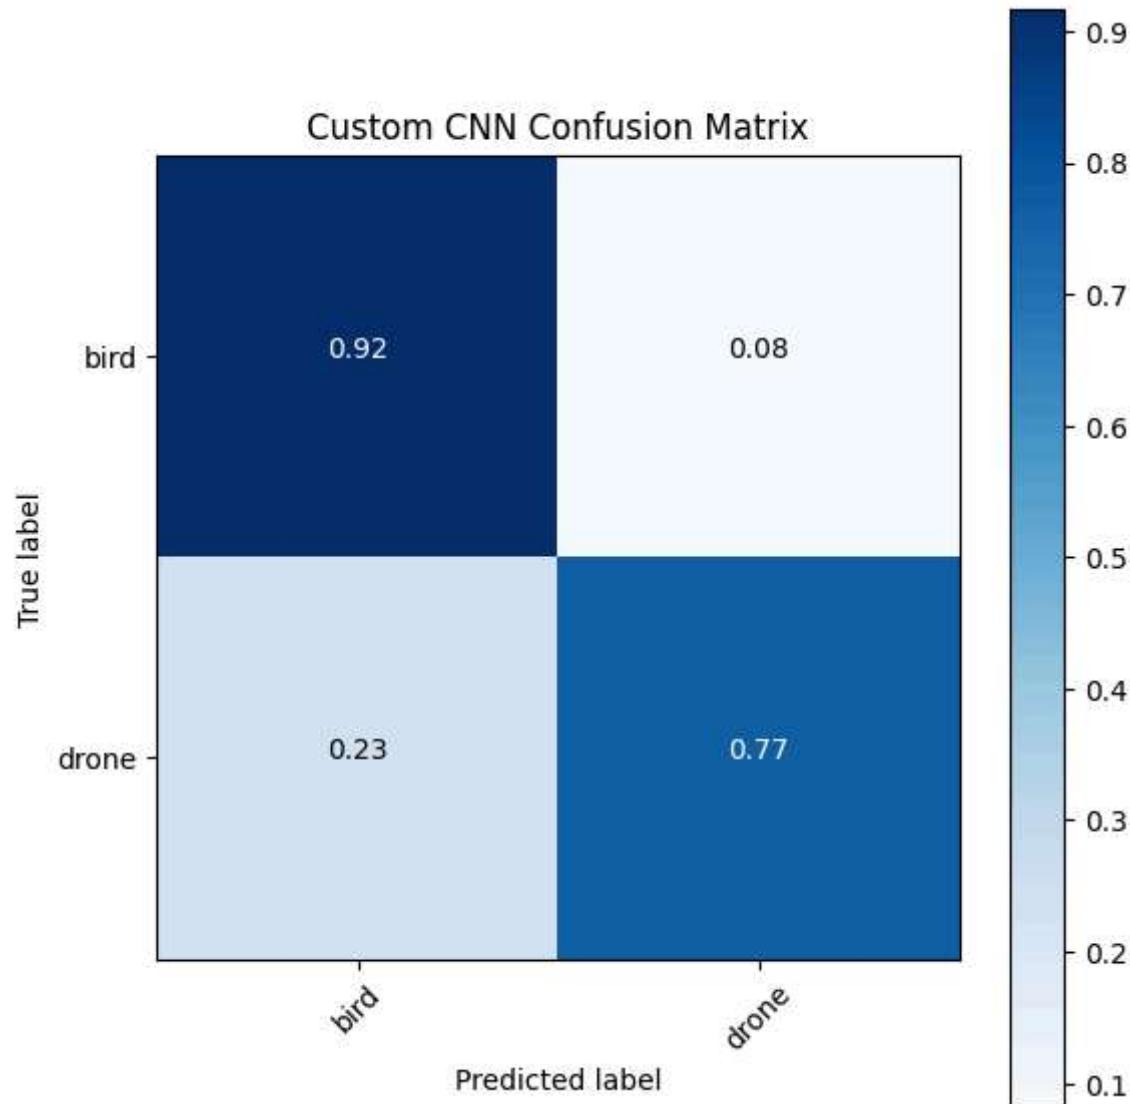
```
7/7 [=====] - 5s 567ms/step
7/7 [=====] - 5s 588ms/step
```

```
In [16]: print("Custom CNN Classification Report:")
print(classification_report(y_true, y_pred_custom, target_names=class_names))

cm_custom = confusion_matrix(y_true, y_pred_custom)
plot_confusion_matrix(cm_custom, classes=class_names, title="Custom CNN Confusion Matrix", normalize=True)
```

Custom CNN Classification Report:

	precision	recall	f1-score	support
bird	0.83	0.92	0.87	121
drone	0.88	0.77	0.82	94
accuracy			0.85	215
macro avg	0.86	0.84	0.85	215
weighted avg	0.85	0.85	0.85	215

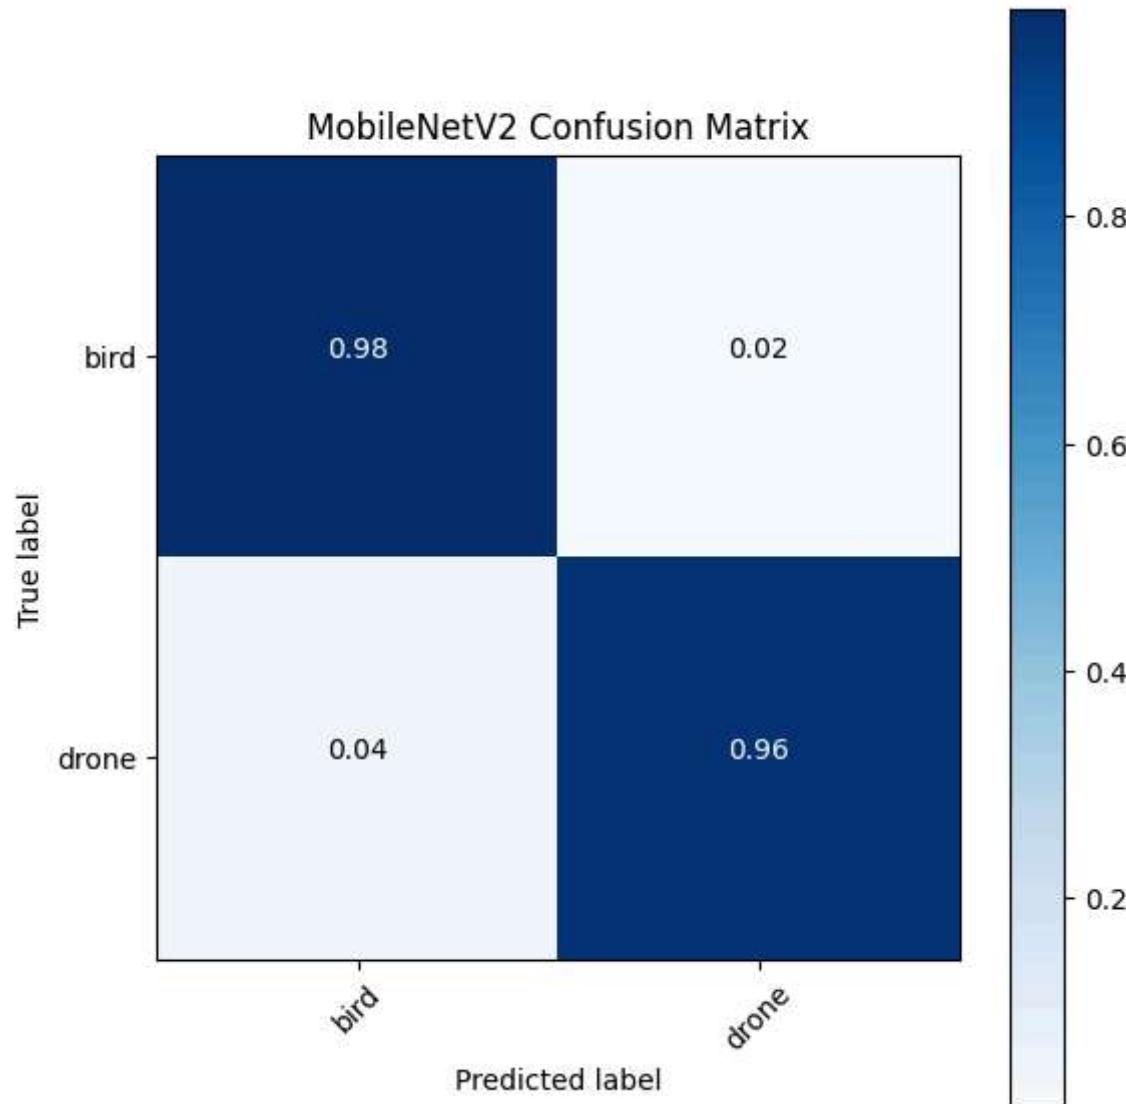


```
In [17]: print("\nMobileNetV2 Classification Report:")
print(classification_report(y_true, y_pred_mobilenet, target_names=class_names))

cm_mobilenet = confusion_matrix(y_true, y_pred_mobilenet)
plot_confusion_matrix(cm_mobilenet, classes=class_names, title="MobileNetV2 Confusion Matrix", normalize=True)
```

MobileNetV2 Classification Report:

	precision	recall	f1-score	support
bird	0.97	0.98	0.98	121
drone	0.98	0.96	0.97	94
accuracy			0.97	215
macro avg	0.97	0.97	0.97	215
weighted avg	0.97	0.97	0.97	215



```
In [18]: from sklearn.metrics import accuracy_score
import pandas as pd
# Compute overall accuracies
acc_custom = accuracy_score(y_true, y_pred_custom)
acc_mobilenet = accuracy_score(y_true, y_pred_mobilenet)

# You can parse classification_report into dicts if you want specific averages:
```

```

report_custom = classification_report(y_true, y_pred_custom, target_names=class_names, output_dict=True)
report_mobilenet = classification_report(y_true, y_pred_mobilenet, target_names=class_names, output_dict=True)

# Macro averages
macro_custom = report_custom["macro avg"]
macro_mobilenet = report_mobilenet["macro avg"]

model_comparison = [
    {
        "Model": "Custom CNN",
        "Test Accuracy": acc_custom,
        "Precision (macro avg)": macro_custom["precision"],
        "Recall (macro avg)": macro_custom["recall"],
        "F1-score (macro avg)": macro_custom["f1-score"],
        "Training Time (min)": custom_train_minutes
    },
    {
        "Model": "MobileNetV2 (Transfer Learning)",
        "Test Accuracy": acc_mobilenet,
        "Precision (macro avg)": macro_mobilenet["precision"],
        "Recall (macro avg)": macro_mobilenet["recall"],
        "F1-score (macro avg)": macro_mobilenet["f1-score"],
        "Training Time (min)": total_mobilenet_minutes
    }
]

df_comparison = pd.DataFrame(model_comparison)
display(df_comparison)

```

	Model	Test Accuracy	Precision (macro avg)	Recall (macro avg)	F1-score (macro avg)	Training Time (min)
0	Custom CNN	0.851163	0.856318	0.841656	0.846099	56.539695
1	MobileNetV2 (Transfer Learning)	0.972093	0.972870	0.970459	0.971576	22.332246

YOLOv8 Object Detection Extension (Bird vs Drone).

In addition to image classification, we prepare the pipeline for **object detection** using **YOLOv8**.

This allows the system not only to classify an image as Bird/Drone but also to **localise** the objects with bounding boxes.

Steps:

1. Install and import YOLOv8 (`ultralytics`).
2. Use the prepared YOLOv8-format dataset (`object_detection_Dataset` with `.txt` labels and `data.yaml`).
3. Train the YOLOv8 model for Bird vs Drone detection.
4. Validate the trained model.
5. Run inference on sample test images.

```
In [19]: # If not already installed (run once in your environment)
# !pip install ultralytics

from ultralytics import YOLO
import os

# Path to YOLOv8 data config (update the path if data.yaml is elsewhere)
DATA_YOLO = r"C:/Users/Shyam/Documents/Labmentix/Project2/Aerial Objection Classification/serialbirddroneProjectV1/datasets/birds_drone_v1"

# Check that file exists
print("YOLO data.yaml exists:", os.path.exists(DATA_YOLO))
print("Using data config:", DATA_YOLO)

# -----
# 1. Load a base YOLOv8 model
# -----
# You can choose yolov8n.pt (nano), yolov8s.pt (small), etc.
yolo_model = YOLO("yolov8n.pt")

# -----
# 2. Train YOLOv8
# -----
# NOTE: This is compute-heavy. Run on GPU if possible.
yolo_model.train(
    data=DATA_YOLO,
    imgsz=640,
    epochs=50,
    batch=16,
```

```
    name="bird_drone_yolov8",
    project="runs/detect"
)
```

YOLO data.yaml exists: True

Using data config: C:/Users/Shyam/Documents/Labmentix/Project2/Aerial Objection Classification/serialbirddroneProject
V1/data/object_detection_Dataset/data.yaml

```
New https://pypi.org/project/ultralytics/8.3.232 available Update with 'pip install -U ultralytics'
Ultralytics YOLOv8.0.196 Python-3.10.10 torch-2.1.0+cpu CPU (12th Gen Intel Core(TM) i7-1255U)
engine\trainer: task=detect, mode=train, model=yolov8n.pt, data=C:/Users/Shyam/Documents/Labmentix/Project2/Aerial Ob
jection Classification/serialbirddroneProjectV1/data/object_detection_Dataset/data.yaml, epochs=50, patience=50, batc
h=16, imgsz=640, save=True, save_period=-1, cache=False, device=None, workers=8, project=runs/detect, name=bird_drone
_yolov8, exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False,
rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, overlap_
mask=True, mask_ratio=4, dropout=0.0, val=True, split=val, save_json=False, save_hybrid=False, conf=None, iou=0.7,
max_det=300, half=False, dnn=False, plots=True, source=None, show=False, save_txt=False, save_conf=False, save_crop=Fa
lse, show_labels=True, show_conf=True, vid_stride=1, stream_buffer=False, line_width=None, visualize=False, augment=Fa
lse, agnostic_nms=False, classes=None, retina_masks=False, boxes=True, format=torchscript, keras=False, optimize=Fals
e, int8=False, dynamic=False, simplify=False, opset=None, workspace=4, nms=False, lr0=0.01, lrf=0.01, momentum=0.937,
weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.
0, kobj=1.0, label_smoothing=0.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5, s
hear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, mosaic=1.0, mixup=0.0, copy_paste=0.0, cfg=None, tracker=botsort.y
aml, save_dir=runs\detect\bird_drone_yolov87
Overriding model.yaml nc=80 with nc=2
```

	from	n	params	module	arguments
0		-1 1	464	ultralytics.nn.modules.conv.Conv	[3, 16, 3, 2]
1		-1 1	4672	ultralytics.nn.modules.conv.Conv	[16, 32, 3, 2]
2		-1 1	7360	ultralytics.nn.modules.block.C2f	[32, 32, 1, True]
3		-1 1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
4		-1 2	49664	ultralytics.nn.modules.block.C2f	[64, 64, 2, True]
5		-1 1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
6		-1 2	197632	ultralytics.nn.modules.block.C2f	[128, 128, 2, True]
7		-1 1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
8		-1 1	460288	ultralytics.nn.modules.block.C2f	[256, 256, 1, True]
9		-1 1	164608	ultralytics.nn.modules.block.SPPF	[256, 256, 5]
10		-1 1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
11		[-1, 6] 1	0	ultralytics.nn.modules.conv.Concat	[1]
12		-1 1	148224	ultralytics.nn.modules.block.C2f	[384, 128, 1]
13		-1 1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
14		[-1, 4] 1	0	ultralytics.nn.modules.conv.Concat	[1]
15		-1 1	37248	ultralytics.nn.modules.block.C2f	[192, 64, 1]
16		-1 1	36992	ultralytics.nn.modules.conv.Conv	[64, 64, 3, 2]
17		[-1, 12] 1	0	ultralytics.nn.modules.conv.Concat	[1]
18		-1 1	123648	ultralytics.nn.modules.block.C2f	[192, 128, 1]
19		-1 1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
20		[-1, 9] 1	0	ultralytics.nn.modules.conv.Concat	[1]
21		-1 1	493056	ultralytics.nn.modules.block.C2f	[384, 256, 1]
22		[15, 18, 21] 1	751702	ultralytics.nn.modules.head.Detect	[2, [64, 128, 256]]

Model summary: 225 layers, 3011238 parameters, 3011222 gradients, 8.2 GFLOPs

Transferred 319/355 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs\detect\bird_drone_yolov87', view at <http://localhost:6006/>
Freezing layer 'model.22.dfl.conv.weight'

train: Scanning C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification\serialbirddroneProjectV1\data

val: Scanning C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification\serialbirddroneProjectV1\data

Plotting labels to runs\detect\bird_drone_yolov87\labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

optimizer: AdamW(lr=0.001667, momentum=0.9) with parameter groups 57 weight(decay=0.0), 64 weight(decay=0.0005), 63 bias(decay=0.0)

Image sizes 640 train, 640 val

Using 0 dataloader workers

Logging results to runs\detect\bird_drone_yolov87

Starting training for 50 epochs...

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
22	1/50	0G	1.443	2.352	1.606	19	640: 100% ██████████ 171/171 [19:26<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
04	2/50	0G	1.586	2.124	1.726	13	640: 100% ██████████ 171/171 [18:10<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
05	3/50	0G	1.658	2.07	1.792	28	640: 100% ██████████ 171/171 [17:59<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
05	4/50	0G	1.649	2.003	1.773	21	640: 100% ██████████ 171/171 [18:00<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
		all	448	663	0.301	0.29	0.189 0.0794

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
08	5/50	0G	1.614	1.918	1.752	23	640: 100% ██████████ 171/171 [17:57<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	6/50	all	448	663	0.32	0.333	0.239 0.108
		0G	1.571	1.786	1.715	37	640: 100% ██████████ 171/171 [17:58<00:00,
04	7/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.542	0.51	0.516 0.237
06	8/50	0G	1.531	1.713	1.682	20	640: 100% ██████████ 171/171 [17:56<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
07	9/50	all	448	663	0.706	0.544	0.625 0.303
		0G	1.504	1.624	1.654	21	640: 100% ██████████ 171/171 [17:54<00:00,
04	10/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.674	0.573	0.614 0.331
04	11/50	0G	1.468	1.578	1.631	29	640: 100% ██████████ 171/171 [17:55<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	11/50	all	448	663	0.513	0.195	0.202 0.101
		0G	1.431	1.511	1.596	13	640: 100% ██████████ 171/171 [17:59<00:00,
04	11/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.651	0.58	0.62 0.336
04	11/50	0G	1.424	1.508	1.597	20	640: 100% ██████████ 171/171 [18:01<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	11/50	all	448	663	0.78	0.625	0.692 0.386

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
05	12/50	0G	1.393	1.462	1.576	18	640: 100% ██████████ 171/171 [17:57<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	13/50	all	448	663	0.781	0.649	0.718 0.399
		0G	1.386	1.44	1.575	20	640: 100% ██████████ 171/171 [17:53<00:00,
04	14/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.693	0.652	0.692 0.403
04	15/50	0G	1.373	1.369	1.564	16	640: 100% ██████████ 171/171 [17:56<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
05	16/50	all	448	663	0.735	0.613	0.676 0.385
		0G	1.355	1.352	1.543	9	640: 100% ██████████ 171/171 [17:57<00:00,
05	17/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.74	0.63	0.696 0.411
05	18/50	0G	1.34	1.297	1.526	33	640: 100% ██████████ 171/171 [17:56<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	18/50	all	448	663	0.83	0.658	0.752 0.446
		0G	1.336	1.297	1.53	16	640: 100% ██████████ 171/171 [17:53<00:00,
04	18/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.786	0.692	0.754 0.433
04	18/50	0G	1.308	1.282	1.522	18	640: 100% ██████████ 171/171 [17:55<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	18/50	all	448	663	0.802	0.647	0.721 0.429

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
04	19/50	0G	1.291	1.219	1.492	16	640: 100% ██████████ 171/171 [17:55<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
05	20/50	all	448	663	0.805	0.669	0.749 0.444
		0G	1.275	1.213	1.497	17	640: 100% ██████████ 171/171 [17:54<00:00,
05	21/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.789	0.692	0.758 0.437
04	22/50	0G	1.254	1.175	1.465	14	640: 100% ██████████ 171/171 [17:52<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	23/50	all	448	663	0.801	0.717	0.766 0.45
		0G	1.262	1.163	1.46	29	640: 100% ██████████ 171/171 [17:58<00:00,
04	24/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.812	0.698	0.762 0.467
05	25/50	0G	1.257	1.145	1.456	15	640: 100% ██████████ 171/171 [17:54<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
05	05	all	448	663	0.795	0.704	0.752 0.469
		0G	1.223	1.112	1.44	21	640: 100% ██████████ 171/171 [17:56<00:00,
05	05	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.796	0.718	0.781 0.462
05	05	0G	1.218	1.115	1.436	25	640: 100% ██████████ 171/171 [17:55<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.857	0.726	0.797 0.479

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
04	26/50	0G	1.213	1.091	1.433	16	640: 100% ██████████ 171/171 [17:53<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
04	27/50	all	448	663	0.823	0.695	0.771 0.478
		0G	1.189	1.065	1.414	14	640: 100% ██████████ 171/171 [17:56<00:00,
04	28/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.766	0.702	0.757 0.47
04	29/50	0G	1.18	1.036	1.412	19	640: 100% ██████████ 171/171 [17:56<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
05	30/50	all	448	663	0.838	0.72	0.788 0.481
		0G	1.169	1.01	1.41	16	640: 100% ██████████ 171/171 [17:51<00:00,
04	31/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.821	0.721	0.792 0.479
04	32/50	0G	1.147	1.002	1.392	41	640: 100% ██████████ 171/171 [17:59<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
06	31/50	all	448	663	0.809	0.701	0.753 0.463
		0G	1.132	0.984	1.379	25	640: 100% ██████████ 171/171 [18:19<00:00,
07	32/50	Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		all	448	663	0.851	0.716	0.792 0.504
07	32/50	0G	1.124	0.9606	1.373	26	640: 100% ██████████ 171/171 [18:39<00:00,
		Class	Images	Instances	Box(P)	R	mAP50 mAP50-95): 100% ██████████ 14/14 [01:
07	32/50	all	448	663	0.877	0.697	0.791 0.48

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
07	33/50	0G	1.137	0.9599	1.381	24	640: 100% ██████████ 171/171 [18:41<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
07	34/50	all	448	663	0.823	0.751	0.804 0.493
		0G	1.104	0.9293	1.342	28	640: 100% ██████████ 171/171 [18:39<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
23	35/50	Class	Images	Instances	Box(P)	R	
		all	448	663	0.873	0.724	0.805 0.507
07	36/50	0G	1.103	0.9383	1.369	23	640: 100% ██████████ 171/171 [19:07<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
07	37/50	all	448	663	0.858	0.738	0.8 0.503
		0G	1.078	0.902	1.342	31	640: 100% ██████████ 171/171 [20:21<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
07	38/50	Class	Images	Instances	Box(P)	R	
		all	448	663	0.866	0.722	0.791 0.484
07	39/50	0G	1.071	0.906	1.337	22	640: 100% ██████████ 171/171 [18:36<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
06	39/50	all	448	663	0.855	0.72	0.798 0.49
		0G	1.069	0.8799	1.32	27	640: 100% ██████████ 171/171 [18:40<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
06	39/50	Class	Images	Instances	Box(P)	R	
		all	448	663	0.858	0.735	0.803 0.505
06	39/50	0G	1.053	0.8827	1.322	23	640: 100% ██████████ 171/171 [18:35<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
06	39/50	all	448	663	0.841	0.755	0.817 0.52

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
06	40/50	0G	1.046	0.8641	1.311	22	640: 100% ██████████ 171/171 [18:40<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.845	0.757 0.81 0.517
							Closing dataloader mosaic
07	41/50	0G	0.9683	0.6788	1.294	8	640: 100% ██████████ 171/171 [18:30<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.869	0.731 0.816 0.513
07	42/50	0G	0.9251	0.6256	1.268	18	640: 100% ██████████ 171/171 [18:29<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.895	0.73 0.825 0.524
14	43/50	0G	0.9062	0.6191	1.251	13	640: 100% ██████████ 171/171 [18:28<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.868	0.745 0.823 0.531
18	44/50	0G	0.8937	0.5938	1.242	27	640: 100% ██████████ 171/171 [20:30<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.874	0.732 0.827 0.522
07	45/50	0G	0.8771	0.5697	1.219	9	640: 100% ██████████ 171/171 [18:24<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	
			all	448	663	0.857	0.745 0.819 0.528
07	46/50	0G	0.8552	0.5557	1.209	11	640: 100% ██████████ 171/171 [18:25<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:
		Class	Images	Instances	Box(P)	R	

		all	448	663	0.91	0.713	0.816	0.522
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
06	47/50	0G Class	0.8306 Images	0.5368 Instances	1.188 Box(P)	8 R	640: 100% ██████████ 171/171 [18:26<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:	
		all	448	663	0.833	0.761	0.818	0.52
07	48/50	0G Class	0.8126 Images	0.5334 Instances	1.184 Box(P)	12 R	640: 100% ██████████ 171/171 [18:26<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:	
		all	448	663	0.866	0.745	0.82	0.525
08	49/50	0G Class	0.8077 Images	0.5242 Instances	1.181 Box(P)	14 R	640: 100% ██████████ 171/171 [18:27<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:	
		all	448	663	0.892	0.739	0.826	0.53
07	50/50	0G Class	0.7983 Images	0.5109 Instances	1.168 Box(P)	9 R	640: 100% ██████████ 171/171 [18:26<00:00, mAP50 mAP50-95): 100% ██████████ 14/14 [01:	
		all	448	663	0.883	0.745	0.824	0.532
50 epochs completed in 16.190 hours.								
Optimizer stripped from runs\detect\bird_drone_yolov87\weights\last.pt, 6.3MB								
Optimizer stripped from runs\detect\bird_drone_yolov87\weights\best.pt, 6.3MB								
Validating runs\detect\bird_drone_yolov87\weights\best.pt...								
Ultralytics YOLOv8.0.196 Python-3.10.10 torch-2.1.0+cpu CPU (12th Gen Intel Core(TM) i7-1255U)								
Model summary (fused): 168 layers, 3006038 parameters, 0 gradients, 8.1 GFLOPs								
58		Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100% ██████████ 14/14 [00:
		all	448	663	0.883	0.743	0.824	0.533
		Bird	448	414	0.857	0.622	0.721	0.417
		Drone	448	249	0.909	0.863	0.927	0.649
Speed: 2.7ms preprocess, 112.8ms inference, 0.0ms loss, 0.6ms postprocess per image								
Results saved to runs\detect\bird_drone_yolov87								

Out[19]: ultralytics.utils.metrics.DetMetrics object with attributes:

```
ap_class_index: array([0, 1])
box: ultralytics.utils.metrics.Metric object
confusion_matrix: <ultralytics.utils.metrics.ConfusionMatrix object at 0x000001267071E650>
fitness: 0.5620609442655043
keys: ['metrics/precision(B)', 'metrics/recall(B)', 'metrics/mAP50(B)', 'metrics/mAP50-95(B)']
maps: array([ 0.41684,  0.64907])
names: {0: 'Bird', 1: 'Drone'}
plot: True
results_dict: {'metrics/precision(B)': 0.882820508623679, 'metrics/recall(B)': 0.7429083522202433, 'metrics/mAP50(B)': 0.8240311895383026, 'metrics/mAP50-95(B)': 0.5329531392351934, 'fitness': 0.5620609442655043}
save_dir: WindowsPath('runs/detect/bird_drone_yolov8t')
speed: {'preprocess': 2.7467720210552216, 'inference': 112.77645719902856, 'loss': 0.0, 'postprocess': 0.6459190377167293}
```

In [22]:

```
# -----
# 3. Validate the trained model
# -----
yolo_model.val()

# Show annotated image
from matplotlib import pyplot as plt
# Ensure inline display
%matplotlib inline
# -----
# 4. Run inference on a sample image
# -----
# Update this path to one image from your test set
test_image_path = "data/object_detection_Dataset/test/images/00ddd713b6f26a17.jpg.rf.0756de13413d353595f8a65e9300b032

if os.path.exists(test_image_path):
    results = yolo_model(test_image_path)
    # Show result in notebook (opens image viewer in some environments)

    # YOLOv8 returns annotated image in BGR → convert to RGB

    annotated_img = results[0].plot()[:, :, ::-1]

    plt.figure(figsize=(10,10))
    plt.imshow(annotated_img)
    plt.axis("off")
```

```
    plt.show()
else:
    print("Test image not found at:", test_image_path)
```

Ultralytics YOLOv8.0.196 Python-3.10.10 torch-2.1.0+cpu CPU (12th Gen Intel Core(TM) i7-1255U)
val: Scanning C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification\serialbirddroneProjectV1\data

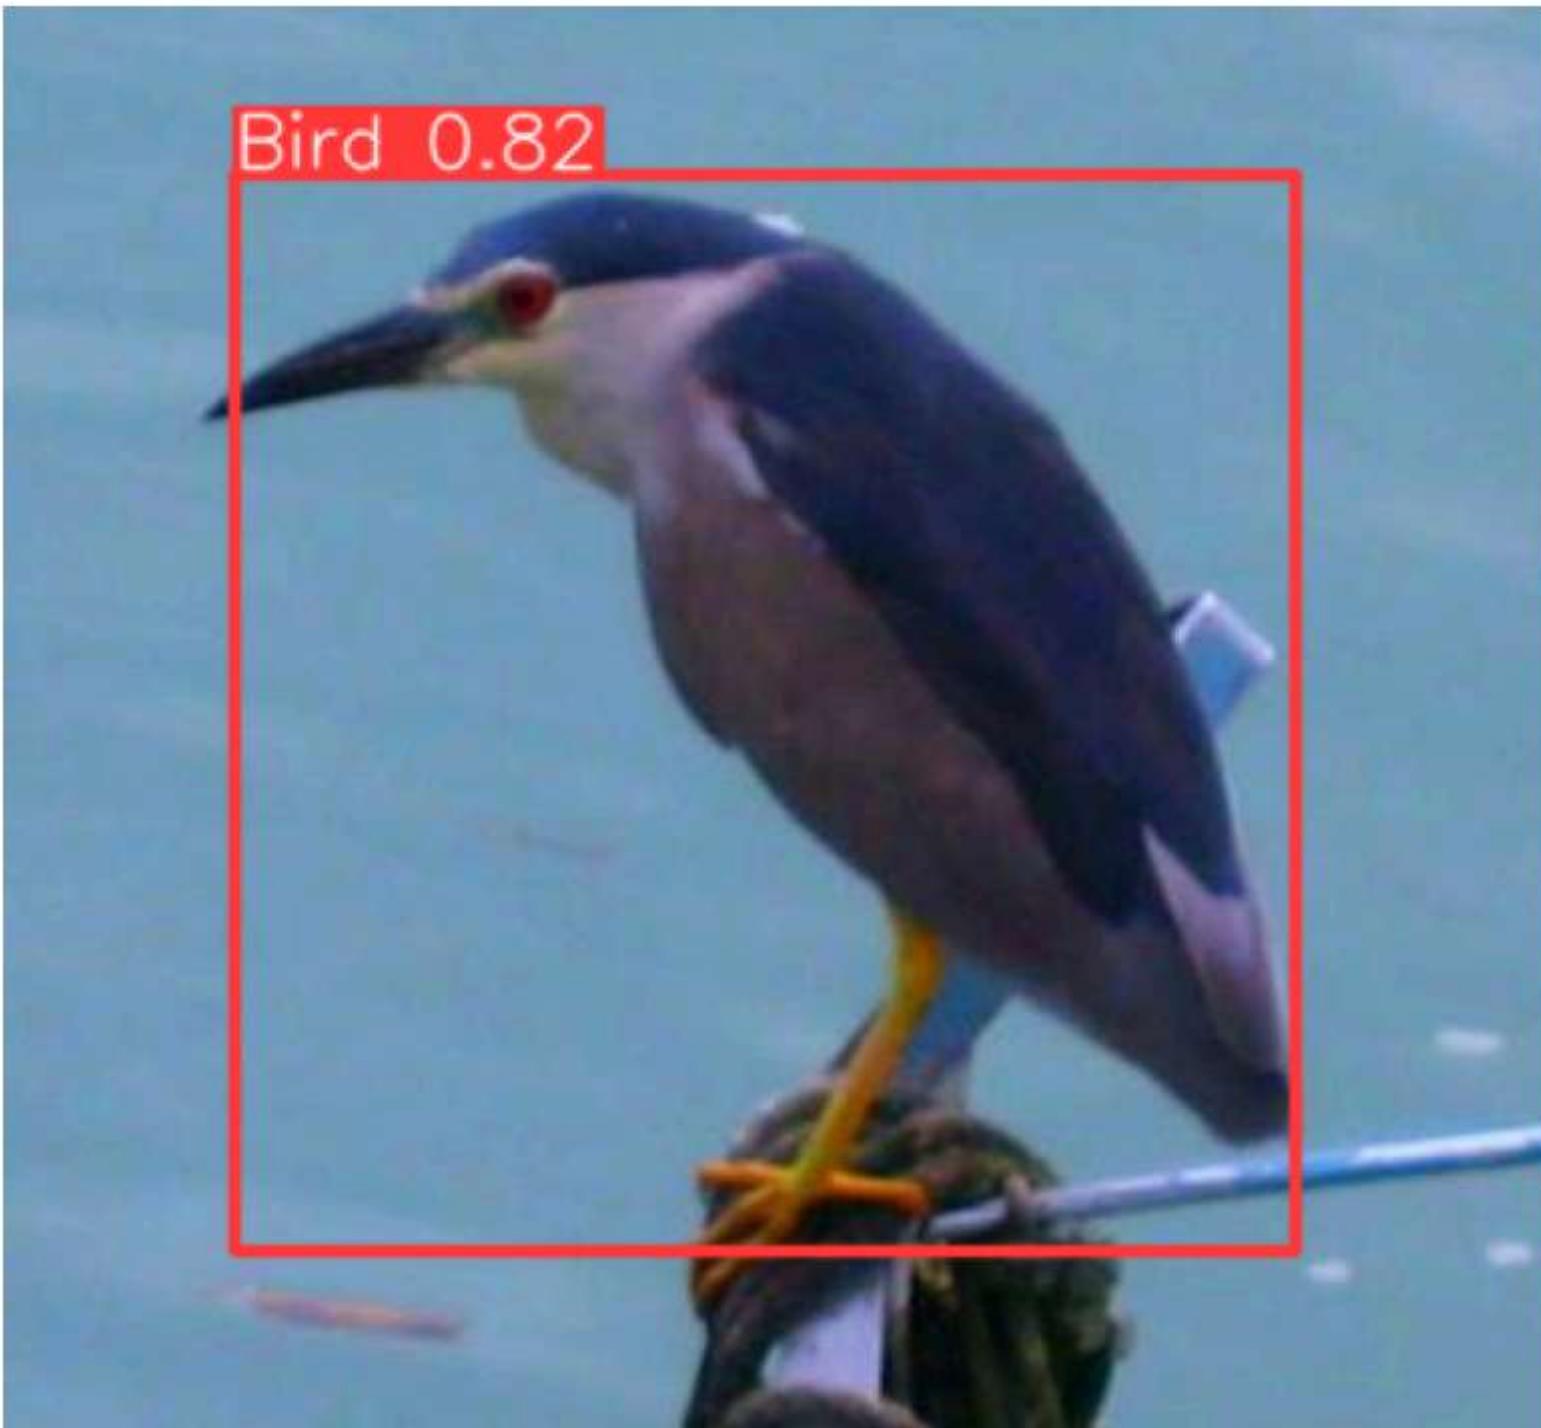
55	Class	Images	Instances	Box(P)	R	mAP50	mAP50-95): 100%	28/28 [00:
	all	448	663	0.883	0.743	0.824	0.533	
	Bird	448	414	0.857	0.622	0.721	0.417	
	Drone	448	249	0.909	0.863	0.927	0.649	

Speed: 2.2ms preprocess, 108.7ms inference, 0.0ms loss, 0.7ms postprocess per image

Results saved to runs\detect\bird_drone_yolov810

image 1/1 C:\Users\Shyam\Documents\Labmentix\Project2\Aerial Objection Classification\serialbirddroneProjectV1\data\object_detection_Dataset\test\images\00ddd713b6f26a17.jpg.rf.0756de13413d353595f8a65e9300b032.jpg: 640x640 1 Bird, 24.9.4ms

Speed: 8.0ms preprocess, 249.4ms inference, 2.0ms postprocess per image at shape (1, 3, 640, 640)





YOLOv8 Results & Notes

- The YOLOv8 model is trained on the same Bird/Drone dataset in **object detection** mode using YOLO annotations.
- This extends the solution from **classification** (single label per image) to **detection** (multiple objects + bounding boxes).
- The trained weights (e.g., `runs/detect/bird_drone_yolov8/weights/best.pt`) are exported and later used in the **Streamlit app** to optionally display detection results with bounding boxes.

Model Evaluation & Final Selection

- Both the **Custom CNN** and **MobileNetV2** models were evaluated on an unseen test set using key performance metrics such as **Accuracy, Precision, Recall, and F1-score**.
- The **fine-tuned MobileNetV2 model** demonstrated superior performance, showing higher recall for the Drone class and a more balanced confusion matrix compared to the Custom CNN.
- Due to its better generalisation capability, faster convergence, and stable validation metrics, the **fine-tuned MobileNetV2** model was selected as the **final model for deployment**.
- This model is used in the Streamlit app to make live predictions for Bird vs Drone classification.

Class Imbalance & Statistical Tests

- The dataset is reasonably balanced between Birds (1752 images) and Drones (1567 images), so no explicit resampling or class-weighting was needed.
- Since this is an image classification problem, model performance was evaluated using classification metrics (Accuracy, Precision, Recall, F1-score) rather than classical statistical hypothesis tests.

Conclusion & Business Recommendations

- The fine-tuned MobileNetV2 model reliably distinguishes between Birds and Drones in aerial images, making it suitable for real-time monitoring applications.
- This solution can be integrated into surveillance systems at airports, wind farms, and restricted zones to automatically flag potential drone intrusions and monitor bird activity.

In []:

In []: